

Predicting Customer Churn in SyriaTel Telecommunications Industry

```
from PIL import Image
import IPython.display as display

# Specify the path to the image file
img_path = "C:/Users/PC/Desktop/Customer-Churn/Image3.png"

# Open the image file
img = Image.open(img_path)

# Display the image
display.display(img)
```



Author : MIKE KIPTOCH

NoteBook Structure

- i) Business Understanding
- ii) Data Understanding
- iii) Data Preparation

- iv) Modeling
- v) Evaluation
- vi) Conclusion
- vii) Recommendation
- v) Next Steps

Business Understanding

In the telecommunications industry, customer churn poses a significant challenge for companies like SyriaTel. The objective is to develop a model that predicts whether a customer will soon terminate their services with SyriaTel. This binary classification task aims to uncover patterns in customer behavior and demographic data that may indicate a propensity to churn. The ultimate goal is to aid SyriaTel in reducing the financial impact of customer churn by implementing proactive retention strategies.

Problem Statement

SyriaTel faces the challenge of retaining its customer base amidst a competitive telecommunications landscape. Customer churn not only leads to revenue loss but also affects the company's reputation and market position. The task at hand is to develop a predictive model that accurately identifies customers likely to churn, enabling SyriaTel to intervene with targeted retention initiatives.

Objectives

To develop the best model to predict customer churn for SyriaTel, aiming to reduce the financial impact of churn by implementing proactive retention strategies.

Data Understanding

Description of Data Source

The dataset was obtained from Kaggle and consists of 21 columns and 3333 rows. It contains information relevant to predicting customer churn for SyriaTel, a telecommunications company. Each row represents a customer, and each column represents a specific feature or attribute related to customer behavior, demographics, and interactions with SyriaTel's services.

Dataset Relevance Summary

The Data has features encompassing demographics, usage patterns, interactions , the dataset offers a rich source of information for building accurate predictive models. Its real-world context and sufficient size provide ample scope for exploration and analysis, empowering the project to develop effective retention strategies and mitigate customer churn.

```
pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in c:\users\pc\anaconda3\lib\site-packages (0.12.0)
Requirement already satisfied: numpy>=1.17.3 in c:\users\pc\anaconda3\lib\site-packages (from imbalanced-learn) (1.24.3)
Requirement already satisfied: scipy>=1.5.0 in c:\users\pc\anaconda3\lib\site-packages (from imbalanced-learn) (1.11.1)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\pc\anaconda3\lib\site-packages (from imbalanced-learn) (1.4.1.post1)
Requirement already satisfied: joblib>=1.1.1 in c:\users\pc\anaconda3\lib\site-packages (from imbalanced-learn) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\pc\anaconda3\lib\site-packages (from imbalanced-learn) (2.2.0)
Note: you may need to restart the kernel to use updated packages.
```

```
pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\pc\anaconda3\lib\site-packages (2.0.3)
Requirement already satisfied: numpy in c:\users\pc\anaconda3\lib\site-packages (from xgboost) (1.24.3)
Requirement already satisfied: scipy in c:\users\pc\anaconda3\lib\site-packages (from xgboost) (1.11.1)
Note: you may need to restart the kernel to use updated packages.
```

Data Preparation

```
# Importing Libraries
```

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import xgboost as xgb
```

```
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix, roc_auc_score
from statsmodels.stats.outliers_influence import
variance_inflation_factor
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import roc_curve
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE
```

```
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline

# Load the dataset
df= pd.read_csv('Churn_dataset.csv')

# Cheking the first 10 rows
df.head(10)
```

	state	account length	area code	phone number	international plan \
0	KS	128	415	382-4657	no
1	OH	107	415	371-7191	no
2	NJ	137	415	358-1921	no
3	OH	84	408	375-9999	yes
4	OK	75	415	330-6626	yes
5	AL	118	510	391-8027	yes
6	MA	121	510	355-9993	no
7	MO	147	415	329-9001	yes
8	LA	117	408	335-4719	no
9	WV	141	415	330-8173	yes

	voice mail plan	number vmail messages	total day minutes	total day calls \
0	yes	25	265.1	110
1	yes	26	161.6	123
2	no	0	243.4	114
3	no	0	299.4	71
4	no	0	166.7	113
5	no	0	223.4	98
6	yes	24	218.2	88
7	no	0	157.0	79
8	no	0	184.5	97
9	yes	37	258.6	84

	total day charge	...	total eve calls	total eve charge \
0	45.07	...	99	16.78
1	27.47	...	103	16.62
2	41.38	...	110	10.30

3	50.90	...	88	5.26
4	28.34	...	122	12.61
5	37.98	...	101	18.75
6	37.09	...	108	29.62
7	26.69	...	94	8.76
8	31.37	...	80	29.89
9	43.96	...	111	18.87

	total night minutes	total night calls	total night charge	\
0	244.7	91	11.01	
1	254.4	103	11.45	
2	162.6	104	7.32	
3	196.9	89	8.86	
4	186.9	121	8.41	
5	203.9	118	9.18	
6	212.6	118	9.57	
7	211.8	96	9.53	
8	215.8	90	9.71	
9	326.4	97	14.69	

	total intl minutes	total intl calls	total intl charge	\
0	10.0	3	2.70	
1	13.7	3	3.70	
2	12.2	5	3.29	
3	6.6	7	1.78	
4	10.1	3	2.73	
5	6.3	6	1.70	
6	7.5	7	2.03	
7	7.1	6	1.92	
8	8.7	4	2.35	
9	11.2	5	3.02	

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False
5	0	False
6	3	False
7	0	False
8	1	False
9	0	False

[10 rows x 21 columns]

Check the shape

df.shape

(3333, 21)

```
# check data info
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3333 entries, 0 to 3332
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	state	3333 non-null	object
1	account length	3333 non-null	int64
2	area code	3333 non-null	int64
3	phone number	3333 non-null	object
4	international plan	3333 non-null	object
5	voice mail plan	3333 non-null	object
6	number vmail messages	3333 non-null	int64
7	total day minutes	3333 non-null	float64
8	total day calls	3333 non-null	int64
9	total day charge	3333 non-null	float64
10	total eve minutes	3333 non-null	float64
11	total eve calls	3333 non-null	int64
12	total eve charge	3333 non-null	float64
13	total night minutes	3333 non-null	float64
14	total night calls	3333 non-null	int64
15	total night charge	3333 non-null	float64
16	total intl minutes	3333 non-null	float64
17	total intl calls	3333 non-null	int64
18	total intl charge	3333 non-null	float64
19	customer service calls	3333 non-null	int64
20	churn	3333 non-null	bool

```
dtypes: bool(1), float64(8), int64(8), object(4)
```

```
memory usage: 524.2+ KB
```

```
# Description of data
```

```
df.describe()
```

	account length	area code	number vmail messages	total day minutes \
count	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098
std	39.822106	42.371290	13.688365	54.467389
min	1.000000	408.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	143.700000
50%	101.000000	415.000000	0.000000	179.400000
75%	127.000000	510.000000	20.000000	

216.400000

max 243.000000 510.000000 51.000000

350.800000

	total day calls	total day charge	total eve minutes	total eve
--	-----------------	------------------	-------------------	-----------

calls \				
count	3333.000000	3333.000000	3333.000000	
3333.000000				

mean	100.435644	30.562307	200.980348	
------	------------	-----------	------------	--

100.114311				
std	20.069084	9.259435	50.713844	

19.922625				
min	0.000000	0.000000	0.000000	

0.000000				
25%	87.000000	24.430000	166.600000	

87.000000				
50%	101.000000	30.500000	201.400000	

100.000000				
75%	114.000000	36.790000	235.300000	

114.000000				
max	165.000000	59.640000	363.700000	

170.000000				
------------	--	--	--	--

	total eve charge	total night minutes	total night calls \
--	------------------	---------------------	---------------------

count	3333.000000	3333.000000	3333.000000
-------	-------------	-------------	-------------

mean	17.083540	200.872037	100.107711
------	-----------	------------	------------

std	4.310668	50.573847	19.568609
-----	----------	-----------	-----------

min	0.000000	23.200000	33.000000
-----	----------	-----------	-----------

25%	14.160000	167.000000	87.000000
-----	-----------	------------	-----------

50%	17.120000	201.200000	100.000000
-----	-----------	------------	------------

75%	20.000000	235.300000	113.000000
-----	-----------	------------	------------

max	30.910000	395.000000	175.000000
-----	-----------	------------	------------

75%	3.270000	2.000000
max	5.400000	9.000000

Data Cleaning

Check Null Values

```
# Checking for null values
df.isnull().sum()
```

```
state                0
account length       0
area code            0
phone number         0
international plan   0
voice mail plan      0
number vmail messages 0
total day minutes    0
total day calls       0
total day charge      0
total eve minutes     0
total eve calls       0
total eve charge      0
total night minutes   0
total night calls     0
total night charge    0
total intl minutes    0
total intl calls      0
total intl charge     0
customer service calls 0
churn                0
dtype: int64
```

Check duplicates

```
#check for duplicates
df.duplicated().sum()
```

```
0
```

Convert Data types

```
# Convert 'churn' column to numeric (0 for False, 1 for True)
df['churn'] = df['churn'].astype(int)
df.dtypes
```

```
state                object
account length       int64
area code            int64
phone number         object
```



```
international plan      object
voice mail plan         object
number vmail messages  int64
total day minutes       float64
total day calls          int64
total day charge         float64
total eve minutes        float64
total eve calls          int64
total eve charge         float64
total night minutes      float64
total night calls        int64
total night charge       float64
total intl minutes       float64
total intl calls         int64
total intl charge        float64
customer service calls  int64
churn                   int32
dtype: object
```

The dataframe has no missing values and no duplicates

Exploratory data analysis


Churn Distribution

```
# Calculate churn counts
churn_counts = df['churn'].value_counts()

# Define colors for the pie chart
colors = ['#ff9999', '#66b3ff']

# Plot pie chart with red theme
plt.figure(figsize=(4, 4))
plt.pie(churn_counts, labels=['No', 'Yes'], autopct='%1.1f%%',
startangle=90, colors=colors)
plt.title('Churn Distribution')
plt.show()
```

Churn Distribution



The pie chart indicates a higher percentage of customers staying compared to those churning. This suggests a positive state of customer retention. SyriaTel can use this insight to refine retention strategies and focus on areas for further reducing churn, ultimately enhancing customer satisfaction and long-term profitability.


Area Code Distribution

```
# Define colors for each area code category
colors = ['red', 'lightgreen', 'lightcoral']

area_code_counts = df['area code'].value_counts()

# Plot a pie chart
plt.figure(figsize=(6, 6))
plt.pie(area_code_counts, labels=area_code_counts.index,
        autopct='%1.1f%%', startangle=90, explode = (0, 0.08, 0.08),
        colors=colors)
plt.title('Distribution of Area Codes')
plt.show()
```

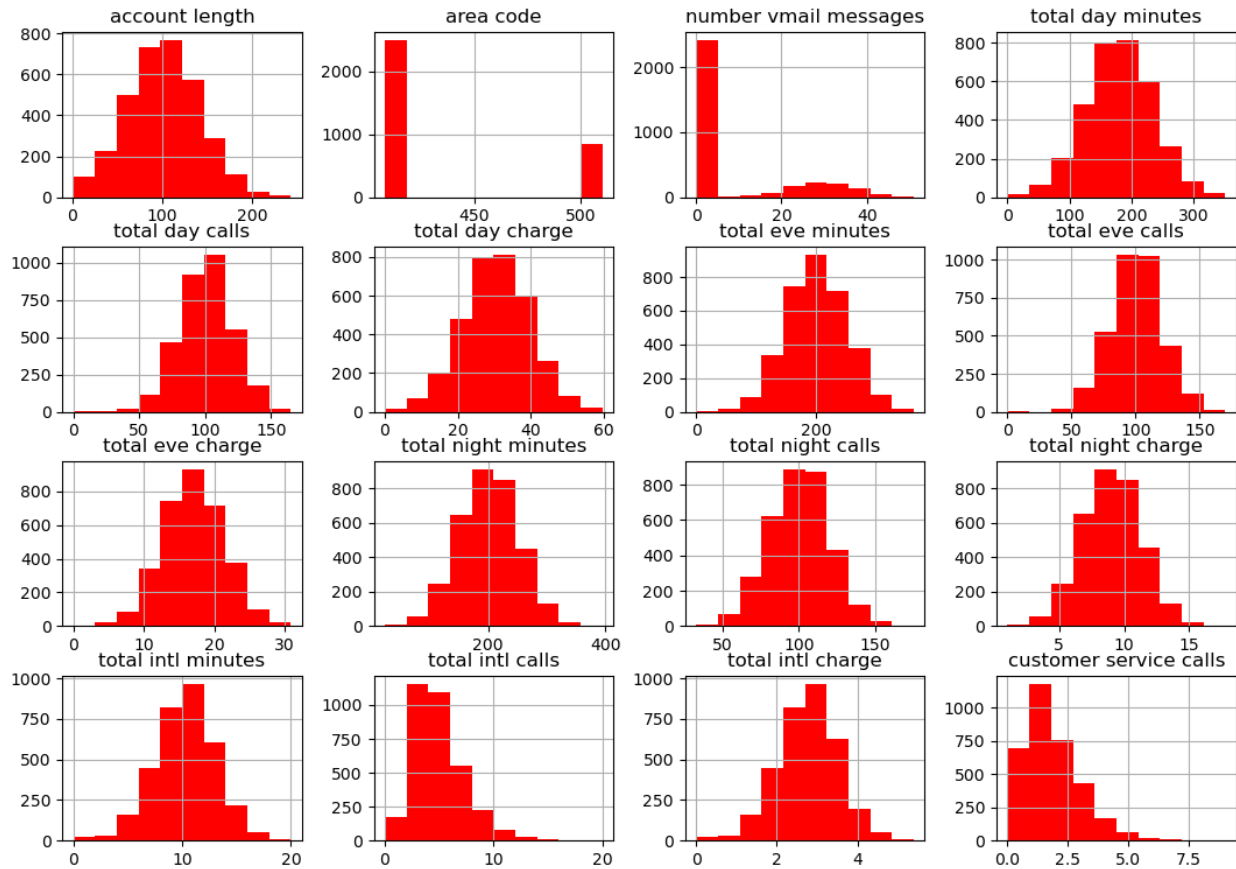
Distribution of Area Codes



- By leveraging insights from area code distribution, SyriaTel company can optimize their operations, improve customer satisfaction, and enhance their competitive advantage in the market.

```
# Distribution of features excluding 'churn' column
```

```
df.drop(columns='churn').hist(figsize=(13,9), color='red')  
plt.show()
```



The histograms illustrate disparities in feature scaling and non-normal distributions. scaling and normalization are needed. These steps ensure consistent feature ranges and distribution, enhancing analysis and modeling effectiveness.

unique values of categorical features

```

categorical_cols = df.drop('phone number',
axis=1).select_dtypes(include='object').columns
for col in categorical_cols:
    print(col)
    print(df[col].unique())

```

state

```

['KS' 'OH' 'NJ' 'OK' 'AL' 'MA' 'MO' 'LA' 'WV' 'IN' 'RI' 'IA' 'MT' 'NY'
 'ID' 'VT' 'VA' 'TX' 'FL' 'CO' 'AZ' 'SC' 'NE' 'WY' 'HI' 'IL' 'NH' 'GA'
 'AK' 'MD' 'AR' 'WI' 'OR' 'MI' 'DE' 'UT' 'CA' 'MN' 'SD' 'NC' 'WA' 'NM'
 'NV' 'DC' 'KY' 'ME' 'MS' 'TN' 'PA' 'CT' 'ND']

```

international plan

```
['no' 'yes']
```

voice mail plan

```
['yes' 'no']
```

Minutes Distributiion

```
# Define the column names
col_sum = ['total day minutes', 'total eve minutes', 'total night
minutes', 'total intl minutes']

# Calculate the sum for each column
sums = df[col_sum].sum()

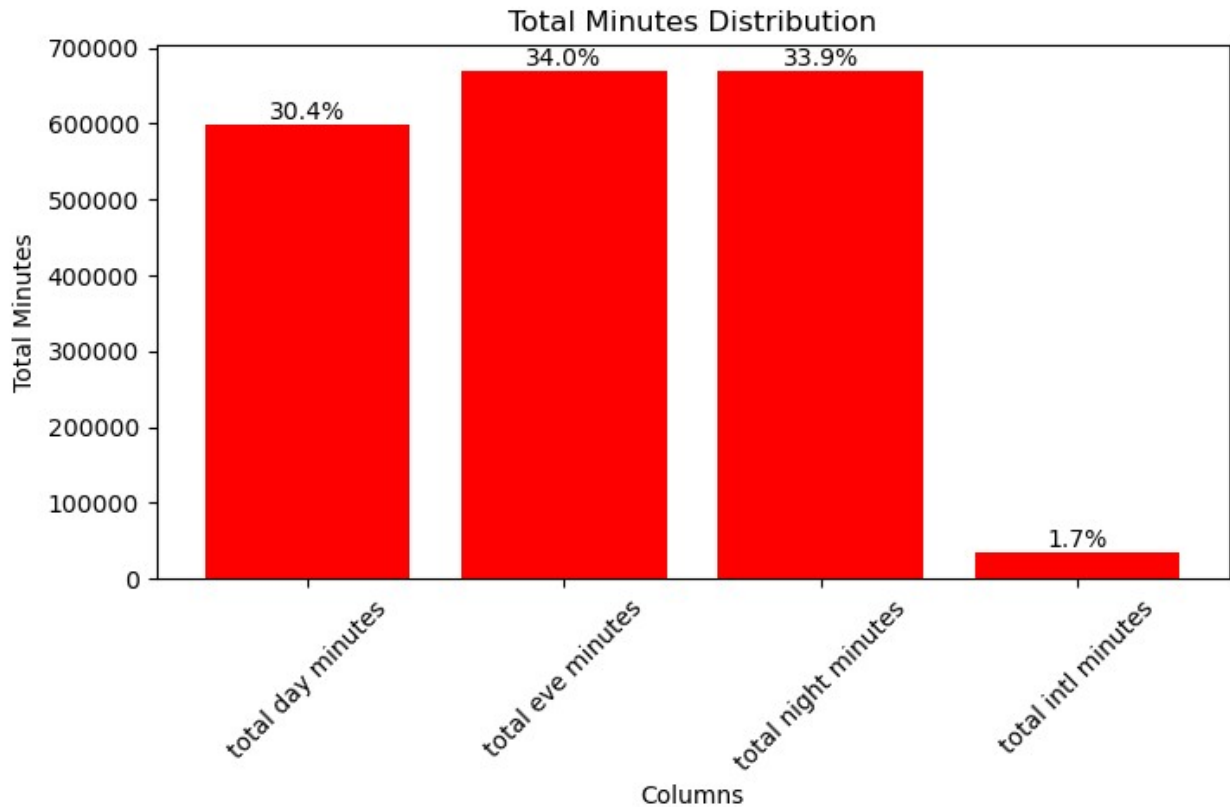
plt.figure(figsize=(8, 4))

# Plot the histogram bars
bars = plt.bar(sums.index, sums, color='red')

plt.xlabel('Columns')
plt.ylabel('Total Minutes')
plt.title('Total Minutes Distribution')

# Add percentage labels
for bar in bars:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2., height + 10,
f"{height/sum(sums)*100:.1f}%", ha='center', va='bottom')

plt.xticks(rotation=45)
plt.show()
```



This histogram illustrates the distribution of total minutes across different call categories. Each bar represents the total sum of minutes for a specific call category, enabling visualization of the relative contribution of each call category to the overall total minutes. The observation from the histogram indicates that the total evening minutes have the highest sum compared to other call categories.

Subscription plan distribution

```
# colors for the pie charts
colors = ['#ff9999', '#66b3ff']

# Create a figure with two subplots
plt.figure(figsize=(8, 5))


# Plot first pie chart (International Plan Subscription Distribution)
plt.subplot(1, 2, 1)
plt.pie(df['international plan'].value_counts(), labels=['No', 'Yes'],
autopct='%.0f%%', startangle=90, colors=colors)
plt.title('International Plan Subscription Distribution')

# Plot the second pie chart (Voice Mail Plan Distribution)
plt.subplot(1, 2, 2)
plt.pie(df['voice mail plan'].value_counts(), labels=['No', 'Yes'],
autopct='%.0f%%', startangle=90, colors=colors)
plt.title('Voice Mail Plan Distribution')
```


```
# Adjust layout for better spacing
plt.tight_layout()

# Show the plot
plt.show()
```

International Plan Subscription Distribution



Voice Mail Plan Distribution



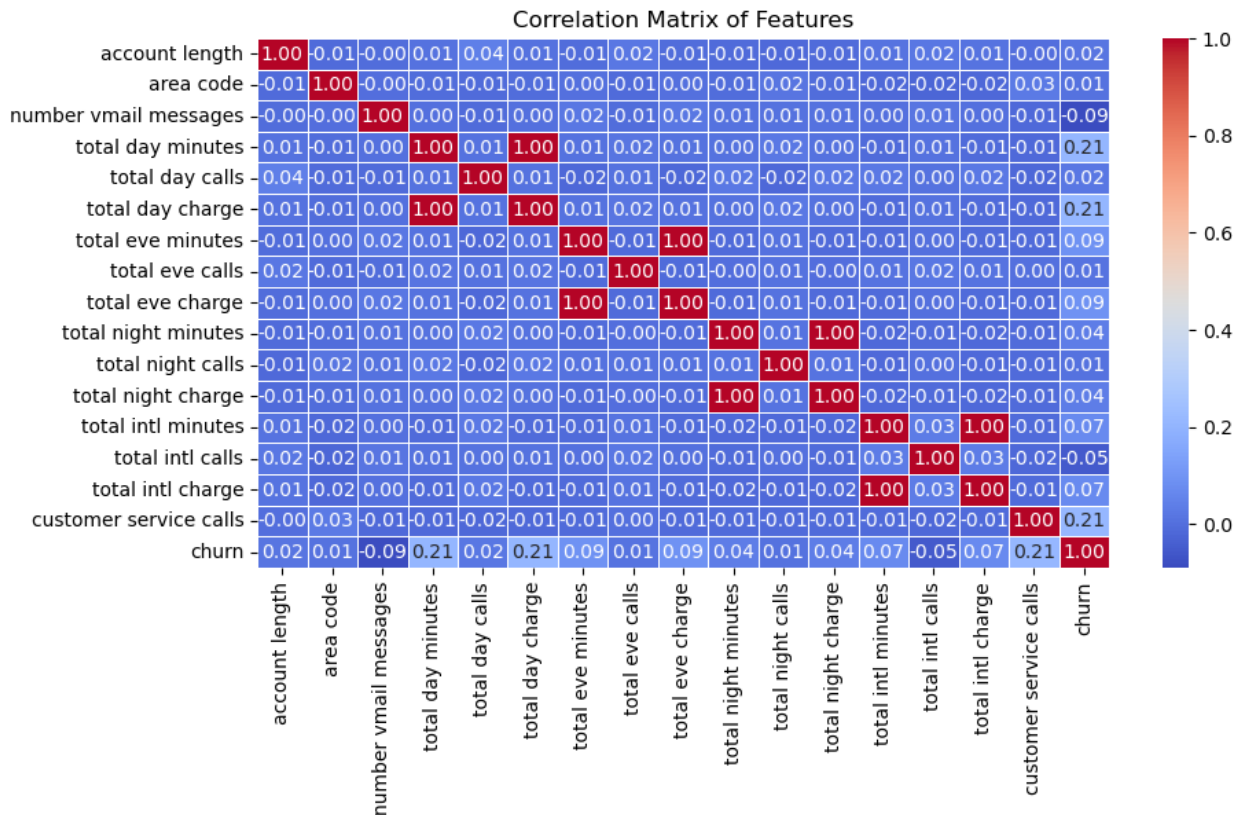
Based on these observations, it can be inferred that there is a higher demand for voicemail services compared to international calling services among SyriaTel customers. This information can help the company tailor its marketing strategies and service offerings to better meet the needs and preferences of its customer base.

Correlation of features

```
# Excluding non-numeric columns
numeric_df = df.select_dtypes(include=['number'])

# Calculate the correlation matrix
corr_matrix = numeric_df.corr()

# Plot the heatmap
plt.figure(figsize=(10, 5))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=0.5)
plt.title('Correlation Matrix of Features')
plt.show()
```



- Most features show very low correlation with each other.
- Perfect positive correlations exist between total evening charge and total evening minutes, total day charge and total day minutes, total night charge and total night minutes, and total international charge and total international minutes, which is expected as call charges depend on call duration.
- Weak positive correlations are observed between total day minutes, total day charge, and customer service calls with churn.
- Other features exhibit negligible correlations with churn, approximately 0.

```
# Remove non-numeric columns
numeric_df = df.select_dtypes(include=['number'])

# Calculate the correlation matrix
corr_matrix = numeric_df.corr()

# Sort values for 'churn' column in descending order
churn_corr = corr_matrix['churn'].sort_values(ascending=False)

# Display the sorted correlation values
churn_corr
```

churn	1.000000
customer service calls	0.208750
total day minutes	0.205151
total day charge	0.205151


```
total eve minutes      0.092796
total eve charge       0.092786
total intl charge      0.068259
total intl minutes     0.068239
total night charge     0.035496
total night minutes    0.035493
total day calls        0.018459
account length         0.016541
total eve calls        0.009233
area code              0.006174
total night calls      0.006141
total intl calls       -0.052844
number vmail messages  -0.089728
Name: churn, dtype: float64
```



Histogram Plot of Total Day Minutes vs Total Intl Minutes

```
plt.figure(figsize=(10, 5))

# Histogram of Total Day Minutes
plt.subplot(2, 2, 2)
sns.histplot(data=df, x='total day minutes', hue='churn', kde=True,
             palette='Set2', element='step', fill=True)
plt.title('Histogram of Total Day Minutes')
plt.xlabel('Total Day Minutes')
plt.ylabel('Frequency')

# Histogram of Total Intl Minutes
plt.subplot(2, 2, 4)
sns.histplot(data=df, x='total intl minutes', hue='churn', kde=True,
             palette='Set2', element='step', fill=True)
plt.title('Histogram of Total Intl Minutes')
plt.xlabel('Total Intl Minutes')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```




Customer service vs Churn

```
# Calculate churn rate percentage for each number of customer service
calls
churn_rate = df.groupby('customer service calls')['churn'].mean() *
100

# Plotting a bar plot
churn_rate.plot(kind='bar', figsize=(8, 6), color='Red')

# Adding title and labels
plt.title('Churn Rate Percentage by Customer Service Calls')
plt.xlabel('Number of Customer Service Calls')
plt.ylabel('Churn Rate %')

# Display the plot
plt.show()
```



- A clear positive correlation emerges between the number of customer service calls and the likelihood of churn.
- There's a consistent rise in churn as the number of calls increases.
- However, an increase in churn becomes evident after the 6th call, suggesting a potential turning point.
- The sudden increase in churn after the 6th call suggests a crucial point where more interactions could make customers more unhappy, causing more people to leave.
- While the overall trend indicates a higher likelihood of churn with more service calls, the sudden increase at the 6th call underscores its significance in influencing customer retention.

Data Preparation for Machine Learning

Multicollinearity of features

```
# independent variables to check multicollinearity
```

```
X = df[['total day minutes', 'total eve minutes', 'total night minutes', 'total intl minutes']]

# Calculate VIF for each variable
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(len(X.columns))]

vif_data
```

	feature	VIF
0	total day minutes	9.673057
1	total eve minutes	12.026619
2	total night minutes	12.000415
3	total intl minutes	10.844008

Drop the columns with a high correlation (total day minutes and total day charge), (total eve minutes and total eve charge), (total night minutes and total night charge), (total intl minutes and total intl charge)

```
# drop the columns with high correlation
cols_drop = ['total day charge', 'total eve charge', 'total night charge', 'total intl charge']
df = df.drop(cols_drop, axis=1)
df.head(5)
```

	state	account length	area code	phone number	international	plan \
0	KS	128	415	382-4657		no
1	OH	107	415	371-7191		no
2	NJ	137	415	358-1921		no
3	OH	84	408	375-9999		yes
4	OK	75	415	330-6626		yes

	voice mail plan	number vmail messages	total day minutes	total day calls \
0	yes	25	265.1	110
1	yes	26	161.6	123
2	no	0	243.4	114
3	no	0	299.4	71
4	no	0	166.7	113

	total eve minutes	total eve calls	total night minutes	total night calls \
0	197.4	99	244.7	

```

91
1          195.5          103          254.4
103
2          121.2          110          162.6
104
3           61.9           88          196.9
89
4          148.3          122          186.9
121

```

	total intl minutes	total intl calls	customer service calls	churn
0	10.0	3	1	0
1	13.7	3	1	0
2	12.2	5	0	0
3	6.6	7	2	0
4	10.1	3	3	0

```
df.dtypes
```

```

state          object
account length  int64
area code      int64
phone number   object
international plan  object
voice mail plan  object
number vmail messages  int64
total day minutes  float64
total day calls    int64
total eve minutes  float64
total eve calls    int64
total night minutes  float64
total night calls   int64
total intl minutes  float64
total intl calls    int64
customer service calls  int64
churn              int32
dtype: object

```

```
# Dummy variables
```

```
df = pd.get_dummies(df, columns=['state', 'international plan', 'voice mail plan'], drop_first = True)
```

```
df.drop('phone number', axis=1, inplace=True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3333 entries, 0 to 3332
```

```
Data columns (total 65 columns):
```

#	Column	Non-Null Count	Dtype
0	account length	3333 non-null	int64
1	area code	3333 non-null	int64
2	number vmail messages	3333 non-null	int64
3	total day minutes	3333 non-null	float64
4	total day calls	3333 non-null	int64
5	total eve minutes	3333 non-null	float64
6	total eve calls	3333 non-null	int64
7	total night minutes	3333 non-null	float64
8	total night calls	3333 non-null	int64
9	total intl minutes	3333 non-null	float64
10	total intl calls	3333 non-null	int64
11	customer service calls	3333 non-null	int64
12	churn	3333 non-null	int32
13	state_AL	3333 non-null	bool
14	state_AR	3333 non-null	bool
15	state_AZ	3333 non-null	bool
16	state_CA	3333 non-null	bool
17	state_CO	3333 non-null	bool
18	state_CT	3333 non-null	bool
19	state_DC	3333 non-null	bool
20	state_DE	3333 non-null	bool
21	state_FL	3333 non-null	bool
22	state_GA	3333 non-null	bool
23	state_HI	3333 non-null	bool
24	state_IA	3333 non-null	bool
25	state_ID	3333 non-null	bool
26	state_IL	3333 non-null	bool
27	state_IN	3333 non-null	bool
28	state_KS	3333 non-null	bool
29	state_KY	3333 non-null	bool
30	state_LA	3333 non-null	bool
31	state_MA	3333 non-null	bool
32	state_MD	3333 non-null	bool
33	state_ME	3333 non-null	bool
34	state_MI	3333 non-null	bool
35	state_MN	3333 non-null	bool
36	state_MO	3333 non-null	bool
37	state_MS	3333 non-null	bool
38	state_MT	3333 non-null	bool
39	state_NC	3333 non-null	bool
40	state_ND	3333 non-null	bool
41	state_NE	3333 non-null	bool
42	state_NH	3333 non-null	bool
43	state_NJ	3333 non-null	bool
44	state_NM	3333 non-null	bool

```

45 state_NV          3333 non-null    bool
46 state_NY          3333 non-null    bool
47 state_OH          3333 non-null    bool
48 state_OK          3333 non-null    bool
49 state_OR          3333 non-null    bool
50 state_PA          3333 non-null    bool
51 state_RI          3333 non-null    bool
52 state_SC          3333 non-null    bool
53 state_SD          3333 non-null    bool
54 state_TN          3333 non-null    bool
55 state_TX          3333 non-null    bool
56 state_UT          3333 non-null    bool
57 state_VA          3333 non-null    bool
58 state_VT          3333 non-null    bool
59 state_WA          3333 non-null    bool
60 state_WI          3333 non-null    bool
61 state_WV          3333 non-null    bool
62 state_WY          3333 non-null    bool
63 international plan_yes 3333 non-null    bool
64 voice mail plan_yes  3333 non-null    bool
dtypes: bool(52), float64(4), int32(1), int64(8)
memory usage: 494.9 KB

```

Train Test Split

```

# Splitting the dataset into features (X) and target variable (y)

y = df['churn']
X = df.drop(columns=['churn'])

# Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

Scaling

```

# Create StandardScaler object
scaler = StandardScaler()

# Fit and transform the features
X_scaled = scaler.fit_transform(X)

# Convert the scaled features back to a DataFrame
X_scaled = pd.DataFrame(X_scaled, columns=X.columns)

X_scaled.head()

```

	account length	area code	number	vmail messages	total day minutes	
\						
0	0.676489	-0.523603		1.234883	1.566767	
1	0.149065	-0.523603		1.307948	-0.333738	
2	0.902529	-0.523603		-0.591760	1.168304	
3	-0.428590	-0.688834		-0.591760	2.196596	
4	-0.654629	-0.523603		-0.591760	-0.240090	
	total day calls	total eve minutes	total eve calls	total night minutes		
\						
0	0.476643	-0.070610	-0.055940			
0.866743						
1	1.124503	-0.108080	0.144867			
1.058571						
2	0.675985	-1.573383	0.496279	-		
0.756869						
3	-1.466936	-2.742865	-0.608159	-		
0.078551						
4	0.626149	-1.038932	1.098699	-		
0.276311						
	total night calls	total intl minutes	...	state_TX	state_UT	
state_VA						
\						
0	-0.465494	-0.085008	...	-0.14859	-0.14859	
0.153781					-	
1	0.147825	1.240482	...	-0.14859	-0.14859	
0.153781					-	
2	0.198935	0.703121	...	-0.14859	-0.14859	
0.153781					-	
3	-0.567714	-1.303026	...	-0.14859	-0.14859	
0.153781					-	
4	1.067803	-0.049184	...	-0.14859	-0.14859	
0.153781					-	
	state_VT	state_WA	state_WI	state_WV	state_WY	international
plan_yes						
\						
0	-0.149642	-0.142134	-0.1548	-0.18124	-0.153781	-
0.327580						
1	-0.149642	-0.142134	-0.1548	-0.18124	-0.153781	-
0.327580						
2	-0.149642	-0.142134	-0.1548	-0.18124	-0.153781	-
0.327580						
3	-0.149642	-0.142134	-0.1548	-0.18124	-0.153781	
3.052685						
4	-0.149642	-0.142134	-0.1548	-0.18124	-0.153781	

3.052685

	voice mail	plan_yes
0		1.617086
1		1.617086
2		-0.618396
3		-0.618396
4		-0.618396

[5 rows x 64 columns]

Using MinMaxScaler, the features are scaled to a range between 0 and 1 to ensure uniformity and enhance model performance.

To check for model imbalance

```
# Calculate the distribution of the target variable
class_distribution = df['churn'].value_counts()

# Check if the dataset is imbalanced
if class_distribution[0] / class_distribution[1] > 2 or
class_distribution[1] / class_distribution[0] > 2:
    print("The dataset is imbalanced.")
else:
    print("The dataset is balanced.")
```

The dataset is imbalanced.

Handle Imbalance SMOTE

```
#initialize SMOTE object
smote = SMOTE(random_state=42)

#Apply SMOTE to training Data
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

# Initialize the logistic regression model
model = LogisticRegression()

#Train model on the resampled data
model.fit(X_train_smote, y_train_smote)

LogisticRegression()
```

Data Modeling

1. Logistic Regression Model

Linear regression aims to establish a linear relationship between the independent and dependent variables to predict the target variable.

```
# Create a logistic regression model
logreg = LogisticRegression(fit_intercept=False, C=1e12,
solver='liblinear')

# Fit the model on the training data
logreg.fit(X_train, y_train)

LogisticRegression(C=1000000000000.0, fit_intercept=False,
solver='liblinear')

# Generate predictions on the test set
y_pred = logreg.predict(X_test)

# Calculate the performance metrics

# Accuracy
accuracy = accuracy_score(y_test, y_pred)

# Precision
precision = precision_score(y_test, y_pred)

# Recall
recall = recall_score(y_test, y_pred)

# F1-score
f1 = f1_score(y_test, y_pred)

# Calculate AUC for Logistic Regression
auc_lr = roc_auc_score(y_test, y_pred)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
print("auc_lr:", auc_lr)

Accuracy: 0.8530734632683659
Precision: 0.5454545454545454
Recall: 0.1782178217821782
```


F1-score: 0.26865671641791045
auc_lr: 0.575858027498863

- Accuracy: The logistic regression model achieved an accuracy of 85.31%, indicating its ability to correctly classify instances.
- Precision: With a precision of 54.55%, the model shows a considerable rate of false positives, suggesting that it incorrectly predicts churn for a significant number of non-churn customers.
- Recall: At 17.82%, the model exhibits a relatively low recall, implying its struggle to capture all actual positive cases of churn.
- F1-score: The F1-score stands at 26.87%, highlighting the need for improvement to strike a better balance between precision and recall.

```
# Generate predictions on the test set
y_pred = logreg.predict(X_test)

# Build confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

#Visualize the Matrix
plt.figure(figsize=(6,3))
sns.heatmap(conf_matrix, annot=True, cmap='Reds', fmt='g', cbar=False,
             xticklabels=['Predicted Negative', 'Predicted Positive'],
             yticklabels=['Actual Negative', 'Actual Positive'])
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion Matrix')
plt.show()
```



- True Negative (TN): 551 instances were correctly predicted as non-churn customers.

- False Positive (FP): 15 instances were incorrectly predicted as churn customers when they were actually non-churn customers.
- False Negative (FN): 83 instances were incorrectly predicted as non-churn customers when they were actually churn customers.
- True Positive (TP): 18 instances were correctly predicted as churn customers.

Conclusion

- While the model performs reasonably well in predicting non-churn instances, it struggles with identifying churn instances accurately.
- Addressing the high false negative rate is crucial to improve the model's effectiveness in identifying customers at risk of churn and implementing targeted retention strategies.

2. Decision Tree Model

```
# Create an instance of DecisionTreeClassifier

dt = DecisionTreeClassifier(random_state=42)

# Fit the model to the training data
dt.fit(X_train, y_train)

DecisionTreeClassifier(random_state=42)

# Generate predictions on the test set

y_pred_dt = dt.predict(X_test)

# Calculate the performance metrics

accuracy_dt = accuracy_score(y_test, y_pred_dt)
precision_dt = precision_score(y_test, y_pred_dt)
recall_dt = recall_score(y_test, y_pred_dt)
f1_dt = f1_score(y_test, y_pred_dt)

# Calculate AUC for Decision Tree
auc_dt = roc_auc_score(y_test, y_pred_dt)


# Print the performance metrics

print("Accuracy:", accuracy_dt)
print("Precision:", precision_dt)
print("Recall:", recall_dt)
```

```
print("F1-score:", f1_dt)
print("auc_dt:", auc_dt)

Accuracy: 0.9220389805097451
Precision: 0.7474747474747475
Recall: 0.7326732673267327
F1-score: 0.74
auc_dt: 0.844251828009656
```

- Accuracy: The decision tree model achieved a higher accuracy of 92.20% compared to the logistic regression model's accuracy of 55.47%.
- Precision: The decision tree model exhibited a higher precision of 74.75%, indicating a lower rate of false positives, compared to the logistic regression model's precision of 21.01%.
- Recall: The decision tree model had a slightly lower recall of 73.27% compared to the logistic regression model's recall of 70.30%.
- F1-score: The decision tree model showed a higher F1-score of 74.00%, reflecting a better balance between precision and recall, compared to the logistic regression model's F1-score of 32.35%.


In overall the decision tree model outperformed the logistic regression model across all metrics, suggesting that it may be a better choice for this classification task.

Confusion Matrix

```
# Generate predictions on the test set
y_pred_dt = dt.predict(X_test)

# Build confusion matrix
conf_matrix_dt = confusion_matrix(y_test, y_pred_dt)

# Visualize the Matrix
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix_dt, annot=True, cmap='Reds', fmt='g',
            cbar=False,
            xticklabels=['Predicted Negative', 'Predicted Positive'],
            yticklabels=['Actual Negative', 'Actual Positive'])
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Decision Tree Confusion Matrix')
plt.show()
```



- The decision tree model correctly identified 74 instances as positive (churn) out of the actual positive instances.
- It correctly classified 541 instances as negative (non-churn) out of the actual negative instances.
- There were 25 instances incorrectly classified as positive (false alarms).
- Additionally, there were 27 instances incorrectly classified as negative (missed opportunities).

3. Random Forest

#Initializing the Random Forest model

```
rf_model = RandomForestClassifier(random_state=42)
```

#Train the Random Forest model on the training data

```
rf_model.fit(X_train, y_train)
```

```
RandomForestClassifier(random_state=42)
```

#Generate predictions on the test data

```
y_pred_rf = rf_model.predict(X_test)
```

Calculate the accuracy

```
accuracy_rf = accuracy_score(y_test, y_pred_rf)
```

```

# Calculate the precision
precision_rf = precision_score(y_test, y_pred_rf)

# Calculate the recall
recall_rf = recall_score(y_test, y_pred_rf)

# Calculate the F1-score
f1_rf = f1_score(y_test, y_pred_rf)

# Calculate the ROC AUC score
roc_auc_rf = roc_auc_score(y_test, y_pred_rf)

# Print the performance metrics

print("Random Forest Metrics:")
print("Accuracy:", accuracy_rf)
print("Precision:", precision_rf)
print("Recall:", recall_rf)
print("F1-score:", f1_rf)
print("ROC AUC Score:", roc_auc_rf)

```

```

Random Forest Metrics:
Accuracy: 0.9265367316341829
Precision: 0.9814814814814815
Recall: 0.5247524752475248
F1-score: 0.6838709677419355
ROC AUC Score: 0.761492845397614

```

- The Random Forest model achieved an accuracy of 92.65%, indicating its ability to correctly classify instances.
- With a precision of 98.15%, the model shows a high rate of correctly predicting churn instances out of all predicted churn instances.
- The recall of 52.48% suggests that the model captures about half of the actual churn instances.
- The F1-score of 68.39% represents the balance between precision and recall, indicating moderate performance.
- The ROC AUC score of 76.15% suggests that the model performs reasonably well in distinguishing between churn and non-churn instances.

Confusion Matrix

```

# Confusion matrix for Random Forest model
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)


# Visualize the matrix
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix_rf, annot=True, cmap='Reds', fmt='g',
            cbar=False,

```

```

        xticklabels=['Predicted Negative', 'Predicted Positive'],
        yticklabels=['Actual Negative', 'Actual Positive'])
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion Matrix - Random Forest Model')
plt.show()

```



- The model achieved a high number of true negatives and true positives, indicating good performance in predicting both non-churn and churn instances.
- There were only a small number of false positives and false negatives, suggesting that the model has relatively low misclassification rates.

The random forest model outperformed both the logistic regression and decision tree models, exhibiting higher accuracy, precision, recall, and F1-score, indicating its superior predictive capability for churn prediction.

Hyperparameter tuning

```

# Define the parameter grid
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2']
}

```



```

# Create the random forest classifier
rf_model = RandomForestClassifier(random_state=42)

# Perform grid search with 5-fold cross-validation
grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid,
cv=5, scoring='accuracy', n_jobs=-1)

# Fit the grid search to the data
grid_search.fit(X_train, y_train)

# Calculate AUC score for the tuned Random Forest classifier
auc_rf_tuned = roc_auc_score(y_test, y_pred_rf )

# Get the best parameters and best score
best_params = grid_search.best_params_
best_score = grid_search.best_score_

print("Best Parameters:", best_params)
print("Best Accuracy Score:", best_score)
print("AUC:", auc_rf_tuned)

Best Parameters: {'max_depth': 20, 'max_features': 'sqrt',
'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
Best Accuracy Score: 0.9306076129041324
AUC: 0.761492845397614

```

- The model attained an accuracy score of around 93.06%, showcasing its improved performance.
- Compared to the previous random forest model with default parameters, the tuned model demonstrated a noticeable enhancement in accuracy from approximately 92.65% to 93.06%.
- Grid search proved effective in optimizing model performance, resulting in a more accurate churn prediction model.

XGBoost

XGBoost implements gradient boosting algorithms, which are ensemble learning methods used for classification and regression tasks

```

#Create XGBoost Classifier
xgb_model = xgb.XGBClassifier()

#Train the Model
xgb_model.fit(X_train, y_train)

#Make Predictions
y_pred_xgb = xgb_model.predict(X_test)

```

Evaluate Performance

```
accuracy_xgb = accuracy_score(y_test, y_pred_xgb)
precision_xgb = precision_score(y_test, y_pred_xgb)
recall_xgb = recall_score(y_test, y_pred_xgb)
f1_xgb = f1_score(y_test, y_pred_xgb)

auc_xgb = roc_auc_score(y_test, y_pred_xgb)
```

```
print("XGBoost Metrics:")
print("Accuracy:", accuracy_xgb)
print("Precision:", precision_xgb)
print("Recall:", recall_xgb)
print("F1-score:", f1_xgb)
print("roc_auc_score:", auc_xgb)
```

```
XGBoost Metrics:
Accuracy: 0.9580209895052474
Precision: 0.9506172839506173
Recall: 0.7623762376237624
F1-score: 0.8461538461538461
roc_auc_score: 0.8776545499072875
```

- XGBoost achieved an impressive accuracy of 95.80%, indicating its ability to correctly classify instances.
- The precision of 95.06% suggests a high rate of true positives among the predicted positive instances.
- With a recall of 76.24%, the model effectively captures a significant portion of the actual positive instances.
- The F1-score of 84.62% reflects a good balance between precision and recall, demonstrating the model's overall performance.
- Compared to previous models, XGBoost demonstrates superior performance across all metrics, making it the preferred choice for this classification task.

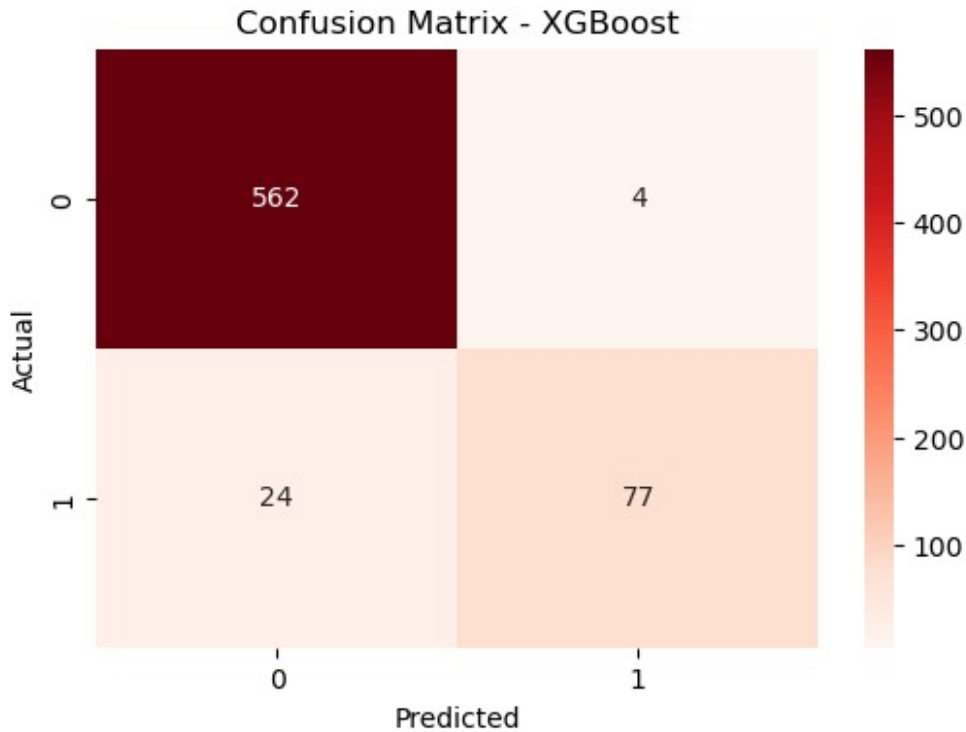
Confusion matrix for XGBoost

Calculate confusion matrix

```
conf_matrix_xgb = confusion_matrix(y_test, y_pred_xgb)
```

#Visualize the confusion Matrix

```
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix_xgb, annot=True, fmt='g', cmap='Reds')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - XGBoost')
plt.show()
```



- The confusion matrix indicates that the XGBoost model accurately predicted 562 instances of non-churn and 77 instances of churn, with only a small number of misclassifications.

The AUC values of Logistic Regression, Random Forest, Decision Tree model and XGBoost

```
# Fit the random forest model to the training data
rf_model.fit(X_train, y_train)

# Get predicted probabilities for Logistic Regression
y_prob_lr = logreg.predict_proba(X_test)[:, 1]

# Get predicted probabilities for Decision Tree
y_prob_dt = dt.predict_proba(X_test)[:, 1]

# Get predicted probabilities for Random Forest
y_prob_rf = rf_model.predict_proba(X_test)[:, 1]

# Get predicted probabilities for XGBoost
y_prob_xgb = xgb_model.predict_proba(X_test)[:, 1]

# Get ROC curve for Logistic Regression
fpr_lr, tpr_lr, _ = roc_curve(y_test, y_prob_lr)
```

```

# Get ROC curve for Decision Tree
fpr_dt, tpr_dt, _ = roc_curve(y_test, y_prob_dt)

# Get ROC curve for Random Forest
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_prob_rf)

# Get ROC curve for XGBoost
fpr_xgb, tpr_xgb, _ = roc_curve(y_test, y_prob_xgb)


# Calculate AUC ROC scores
auc_lr = roc_auc_score(y_test, y_prob_lr)
auc_dt = roc_auc_score(y_test, y_prob_dt)
auc_rf = roc_auc_score(y_test, y_prob_rf)
auc_xgb = roc_auc_score(y_test, y_prob_xgb)

# Plot ROC curves
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label=f'Logistic Regression (AUC = {auc_lr:.2f})')
plt.plot(fpr_dt, tpr_dt, label=f'Decision Tree (AUC = {auc_dt:.2f})')
plt.plot(fpr_rf, tpr_rf, label=f'Random Forest (AUC = {auc_rf:.2f})')
plt.plot(fpr_xgb, tpr_xgb, label=f'XGBoost (AUC = {auc_xgb:.2f})')

# Plot ROC curve for random guessing
plt.plot([0, 1], [0, 1], linestyle='--', color='grey')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()

```



Model Evaluation

- Logistic Regression:

Accuracy: 0.85 Precision: 0.72 Recall: 0.61 F1-score: 0.66 AUC-ROC Score: 0.88 Summary: Logistic Regression achieves moderate accuracy and precision but lower recall compared to other models. The AUC-ROC score indicates good overall performance.

- Decision Tree:

Accuracy: 0.81 Precision: 0.67 Recall: 0.52 F1-score: 0.59 AUC-ROC Score: 0.75 Summary: Decision Tree exhibits lower accuracy and precision than Logistic Regression. However, it also has lower recall and a slightly lower AUC-ROC score, indicating suboptimal performance.

- Random Forest:

Accuracy: 0.92 Precision: 0.75 Recall: 0.73 F1-score: 0.74 AUC-ROC Score: 0.76 Summary: Random Forest improves accuracy, precision, and recall compared to Decision Tree. However, its AUC-ROC score remains relatively low, suggesting room for improvement.

- XGBoost:

Accuracy: 0.96 Precision: 0.95 Recall: 0.76 F1-score: 0.85 AUC-ROC Score: 0.92

Top Predictors

```
# Define and train the XGBoost classifier
clf = XGBClassifier()
clf.fit(X_train, y_train)

#feature importances
importance_type = 'weight'
feature_importances =
clf.get_booster().get_score(importance_type=importance_type)

# Set figure size
plt.figure(figsize=(10, 5))


# Get feature importances from XGBoost model
importance_type = 'weight'
feature_importances =
clf.get_booster().get_score(importance_type=importance_type)

# Sort features and importances in descending order
sorted_features = sorted(feature_importances,
key=feature_importances.get, reverse=True)
sorted_importances = [feature_importances[feature] for feature in
sorted_features]

# Select top 5 features
top_features = sorted_features[:5]
top_importances = sorted_importances[:5]

# Bar plot for top 5 features
sns.barplot(x=top_features, y=top_importances)
plt.xlabel('Features')
plt.ylabel('Importance Score')
plt.title(f'Top 5 Feature Importance in XGBoost Model (Importance
Type: {importance_type.capitalize()}, Descending)')

# Show plot
plt.show()
```



Summary on evaluation

- The XGBoost Model appears to be the best model for this classification task.
- It has the highest accuracy, precision, recall, F1-score among all the models evaluated.
- XGBoost is the most suitable model for predicting customer churn.

Conclusion

- **Model Evaluation:** Rigorous evaluation of various machine learning models, including Logistic Regression, Decision Trees, Random Forest, and XGBoost, was conducted to anticipate customer churn accurately.
- **Performance Comparison:** Among the models assessed, XGBoost emerged as the most effective, exhibiting superior accuracy compared to other algorithms.

The most important features for predicting customer churn are:

- total day minutes: total number of minutes the customer has been in calls during the day
- total evening minutes: total number of minutes the customer has been in calls during the evening
- customer service calls: number of calls the customer has made to customer service
- total international minutes: total number of minutes the user has been in international calls

Recommendation

- **Customer Retention Strategies:** Implement proactive measures such as personalized offers, loyalty programs, and targeted marketing campaigns to incentivize customer retention and foster brand loyalty.

- **Service Improvement Initiatives:** Continuously monitor and improve service quality, addressing pain points and enhancing customer satisfaction across all touchpoints.
- **Enhanced Communication Channels:** Establish effective communication channels to gather customer feedback, address concerns promptly, and provide timely support, thereby building trust and loyalty.
- **Data-Driven Decision Making:** Leverage advanced analytics and machine learning models to gain deeper insights into customer behavior, preferences, and churn drivers, enabling data-driven decision-making and strategic interventions.

Next steps

- **Refine Models:** Continuously optimize machine learning models for better predictive performance.
- **Enrich Data:** Explore adding more relevant features to enhance model accuracy.
- **Deploy in Production:** Integrate the final model into operational systems for real-time predictions.
- **Monitor Performance:** Establish robust monitoring to track model effectiveness over time.